

SDE-2

Overview

As an **SDE-2**, your task is to build a production-grade data pipeline that ingests, transforms, validates, and stores agricultural sensor data coming from multiple sources. The goal is to deliver clean, enriched, and query-optimized data for downstream analytics.

Our agricultural monitoring system collects data from sensors placed across farmlands. These sensors track parameters such as **soil moisture**, **temperature**, **humidity**, **light intensity**, and **battery levels**. The raw data is stored in Parquet files, split by date and sensor type. Your job is to build a pipeline to process, validate, and enrich this data.

Tasks

1. Data Ingestion

Build a modular ingestion component that:

- Reads daily Parquet files from a `data/raw/` directory (local or mounted cloud storage).

Example file: `data/raw/2023-06-01.parquet`

(<https://drive.google.com/file/d/1JzvmQU1ETr4MBg0zTYbpjEm43QV0wXIW/view?usp=sharing>)

Schema:

```
sensor_id (string)
timestamp (ISO datetime)
reading_type (string) -- e.g., temperature, humidity
value (float)
battery_level (float)
```

- Supports **incremental loading** (e.g., using file naming or timestamp-based checkpointing).
- Uses **DuckDB** to:
 - Inspect the file schema quickly.
 - Run validation queries (e.g., missing columns, data ranges).
 - Log ingestion summary via DuckDB SQL aggregation.
- Handles and logs:
 - Corrupt or unreadable files
 - Schema mismatches
 - Missing or invalid values
- **Logs ingestion statistics:**
 - Number of files read

- Records processed
- Records skipped/failed

2. Data Transformation

Apply the following transformation logic:

- **Data cleaning:**
 - Drop duplicates
 - Fill or drop missing values
 - Detect and correct outliers (e.g., z-score > 3)
- **Derived fields:**
 - Daily average value per sensor and reading_type
 - 7-day rolling average per sensor
 - `anomalous_reading = true` if value outside expected range (define a simple range per reading_type)
 - Normalize using calibration logic (e.g., `value = raw_value * multiplier + offset` — hardcode/sample calibration params)
- **Timestamp processing:**
 - Convert to ISO 8601 format
 - Adjust to UTC+5:30

3. Data Quality Validation (Using DuckDB)

Use **DuckDB queries and/or validation framework** to:

- Validate types (e.g., value is float, timestamp is ISO format)
- Check expected value ranges per `reading_type`
Detect **gaps in hourly data** (use DuckDB `generate_series` to simulate expected times)
- Profile:
 - % missing values per reading_type
 - % anomalous readings
 - Time coverage gaps per sensor

Save the validation output as `data_quality_report.csv`

4. Data Loading & Storage

- Store cleaned and transformed data in **Parquet** format under `data/processed/`
- Optimize for analytical queries:
 - Columnar format
 - Partition by `date` and optionally by `sensor_id`
 - Apply compression (e.g., `snappy`, `gzip`)

Deliverables

1. **Codebase** in **any language of your choice** that implements the pipeline described above:
 - Modular structure for ingestion, transformation, validation, and loading
 - Code should be clean, documented, and testable
2. **Unit tests** for core logic (e.g., transformation, anomaly detection, validation)
3. **Unit test coverage**
4. **Docker setup:**
 - Dockerfile to build and run your code locally
5. **Sample Data:**
 - Include at least 3 Parquet files in `data/raw/` (e.g., 3 different dates) to demonstrate your pipeline
6. **Documentation:**
 - README with:
 - Setup & run instructions
 - Overview of architecture and components
 - Explanation of calibration & anomaly logic
 - Example of generated **data quality report**
7. **Submission:**
 - Push your code and sample data to a **public GitHub repository**
 - Include a short demo video
 - Ingestion
 - Transformation
 - Validation
 - Sample query on DuckDB

For any doubts or queries, reach out to us at : aravindhan@satsure.co